


Virtual Me (VME)

The AI-Native Digital Twin
and Multi-Agent Workspace



Bridging the Context Chasm


Virtual Me (VME) operates as an AI-native digital twin. It permanently aligns human long-term memory with a persistent, dynamic LLM context environment—acting as the ultimate memory bank for complex development.



The Active Developer Workspace

Time-Travel Debugging

Effortlessly reload historical context from a year ago to root-cause legacy bugs introduced at specific commits.



Virtual Teams

Instantiate distinct personas (like Alice & Bob) that dynamically merge specialized context for complex problem-solving.


Skill Manager

Inject reusable prompts, architectural rules, and system instructions directly into the AI environment.

Tech Stack

Distraction-free minimalism powered by lightning-fast React 19, Vite, and Tailwind CSS v4.


Multi-Agent Orchestration Engine




Low-Latency SSE Tech

Utilizes X-Accel-Buffering: no and Connection: keep-alive headers. This prevents reverse-proxy buffering, delivering real-time token streaming and agent thought-processes directly to the UI without delays.

Human-in-the-Loop & Autonomous Evolution




Cloud-Native Scale & Smart Context Economics



The Economic Breakthrough

Maxing out a 2M token window per query is cost-prohibitive. VME relies on dynamic RAG routing to pull only relevant history into a cheaper 200K hot token window.



We aren't just scaling up 1000x; we are providing virtually unlimited, cost-effective **smart context**.